

# Tech-Triathlon Datathon Challenge

This document outlines the comprehensive data preprocessing pipeline implemented for the government service prediction system. Our approach transforms raw booking, task, and staffing data into machine-learning-ready features for two primary prediction tasks: service completion time estimation and staffing requirement forecasting.



## Input Datasets

- **Bookings Dataset:** 11 columns containing citizen appointment details and processing times
- **Tasks Dataset:** 4 columns mapping task IDs to sections (requires manual completion)
- **Staffing Dataset:** 4 columns with daily staffing levels and workload metrics

## Data Quality Assessment

- **Missing Values:** Minimal missing data in core datasets
- **Data Types:** Mixed numeric, categorical, and datetime fields
- **Outliers:** Processing times > 8 hours removed as anomalous
- **Temporal Range:** Historical data from 2021-2024

## Phase 1: Tasks Dataset Completion

### Challenge

The tasks dataset arrived with empty `task_name` and `section_name` fields, requiring manual completion while preserving existing IDs.

### Solution

```
# Define 6 government sections
section_mapping = {
    'SEC-001': 'First-time Passport Applications',
    'SEC-002': 'Renewals and Updates',
    'SEC-003': 'Corrections and Amendments',
    'SEC-004': 'Lost/Stolen Passport Reissue',
    'SEC-005': 'Document Verification',
    'SEC-006': 'Special Cases'
}
```

### Rationale

- **Domain Alignment:** Chose Immigration & Emigration services for realistic government context
- **Balanced Distribution:** Ensured logical task distribution across sections
- **ID Preservation:** Maintained original task\_id and section\_id mappings

## Phase 2: Target Variable Engineering

### Task 1: Processing Time Calculation

```
processing_time_minutes = (check_out_time - check_in_time).total_seconds()
/ 60
```

### Data Cleaning Steps:

- Removed negative processing times (data entry errors)
- Capped maximum processing time at 480 minutes (8 hours)

- Filtered out extreme outliers using IQR method

### Statistical Summary:

- Mean processing time: ~48 minutes
- Standard deviation: ~24 minutes
- Range: 5-217 minutes after cleaning

### Task 2: Employee Count Target

- Direct extraction from `employees_on_duty` field
- No transformation required as already in target format
- Range: 1-8 employees per section per day

## Phase 3: Feature Engineering

### Temporal Features

#### Date/Time Decomposition:

```
# Extract meaningful time components
day_of_week = appointment_date.dt.dayofweek
month = appointment_date.dt.month
hour = appointment_time.dt.hour
is_weekend = day_of_week.isin([5, 6])
```

#### Rationale:

- **Day of Week:** Government offices have different patterns on weekdays vs weekends
- **Hour:** Processing times vary by appointment time (morning rush, lunch breaks)
- **Month:** Seasonal variations in service demand
- **Weekend Flag:** Binary feature for non-working days

### Categorical Encoding

#### Label Encoding Strategy:

```
# Fit encoders on combined train+test data
all_task_ids = set(train_tasks) | set(test_tasks)
le_task.fit(all_task_ids)
```

#### Benefits:

- Handles unseen categories in test data
- Preserves ordinal relationships where applicable
- Memory efficient compared to one-hot encoding

### Historical Aggregation Features

## Task-Level Averages:

```
task_avg_time = bookings.groupby('task_id')['processing_time'].mean()
section_avg_time = bookings.groupby('section_id')['processing_time'].mean()
```

## Feature Types:

- **Task Average Time:** Historical mean processing time per task type
- **Section Average Time:** Historical mean processing time per section
- **Hour Average Time:** Time-of-day patterns in processing duration
- **Section Average Employees:** Historical staffing patterns
- **Workload Metrics:** Total task time per employee ratios

## Rationale:

- Captures domain-specific patterns not evident in raw features
- Provides baseline estimates for new/unseen combinations
- Reduces model complexity by pre-computing statistical patterns

# Phase 4: Data Integration and Validation

## Dataset Joining Strategy

```
# Join bookings with task information
enhanced_df = bookings.merge(tasks, on='task_id', how='left')
```

## Validation Checks:

- Verified all task\_ids have corresponding section mappings
- Ensured no data leakage between training and test sets
- Confirmed temporal consistency across datasets

## Missing Value Handling

### Strategy by Column:

- **Categorical:** Mode imputation or "Unknown" category
- **Numerical:** Mean/median imputation based on distribution
- **Temporal:** Forward-fill for time series patterns
- **Historical Averages:** Global mean when specific patterns unavailable

# Phase 5: Feature Selection and Scaling

## Final Feature Sets

### Task 1 (Processing Time):

- Temporal: day\_of\_week, month, hour, is\_weekend

- Categorical: task\_id\_encoded, section\_id\_encoded
- Contextual: num\_documents, queue\_number
- Historical: task\_avg\_time, section\_avg\_time, hour\_avg\_time

### Task 2 (Staffing Needs):

- Temporal: day\_of\_week, month, is\_weekend
- Categorical: section\_id\_encoded
- Workload: total\_task\_time\_minutes, section\_avg\_workload
- Historical: section\_avg\_employees

## Scaling Decision

**No explicit scaling applied** for Random Forest models as they are:

- Tree-based algorithms (scale-invariant)
- Handle mixed data types naturally
- Robust to outliers without preprocessing

## Results and Validation

### Data Quality Metrics

- **Completeness:** 99.8% complete after preprocessing
- **Consistency:** All temporal relationships validated
- **Coverage:** Test set categories covered in training data

### Feature Importance Analysis

#### Top contributors for Task 1:

1. Historical task averages (81% importance)
2. Queue number (4.9% importance)
3. Month (3.5% importance)

#### Top contributors for Task 2:

1. Total task time minutes (98% importance)
2. Month (0.8% importance)

## Challenges and Solutions

### Challenge 1: Unseen Categories in Test Data

**Problem:** Test data contained task\_ids not in training set **Solution:** Implemented safe encoding with fallback defaults

### Challenge 2: Temporal Consistency

**Problem:** Future prediction dates beyond training range **Solution:** Extracted cyclical features (day, month) rather than absolute dates

## Challenge 3: Missing Historical Context

**Problem:** New tasks/sections without historical averages **Solution:** Global mean imputation with domain-reasonable defaults

## Conclusion

The preprocessing pipeline successfully transformed raw government service data into robust machine-learning features. Key innovations include:

- **Comprehensive temporal decomposition** capturing government office patterns
- **Multi-level historical aggregation** providing context-aware baselines
- **Robust encoding strategies** handling unseen test scenarios
- **Domain-informed feature engineering** leveraging government service knowledge

This foundation enables accurate prediction of both service completion times and staffing requirements, directly supporting the datathon's goal of optimizing citizen service delivery.